

Minsky Cluster Training

—
Eric Fiala
Cognitive Solution Architect

Agenda

Cluster Basics

Cluster Overview
Cluster Access

Spark

Spark Instance Group Overview
Submitting Basic Spark Jobs

Deep Learning

Deep Learning Overview
Importing Datasets
Running Basic Training Jobs
Hyperparameter Tuning
Elastic Distributed Training

Jupyter / Anaconda

Jupyter / Anaconda overview
Work with anaconda through notebook
Work with anaconda through CLI

Open Discussion / Questions

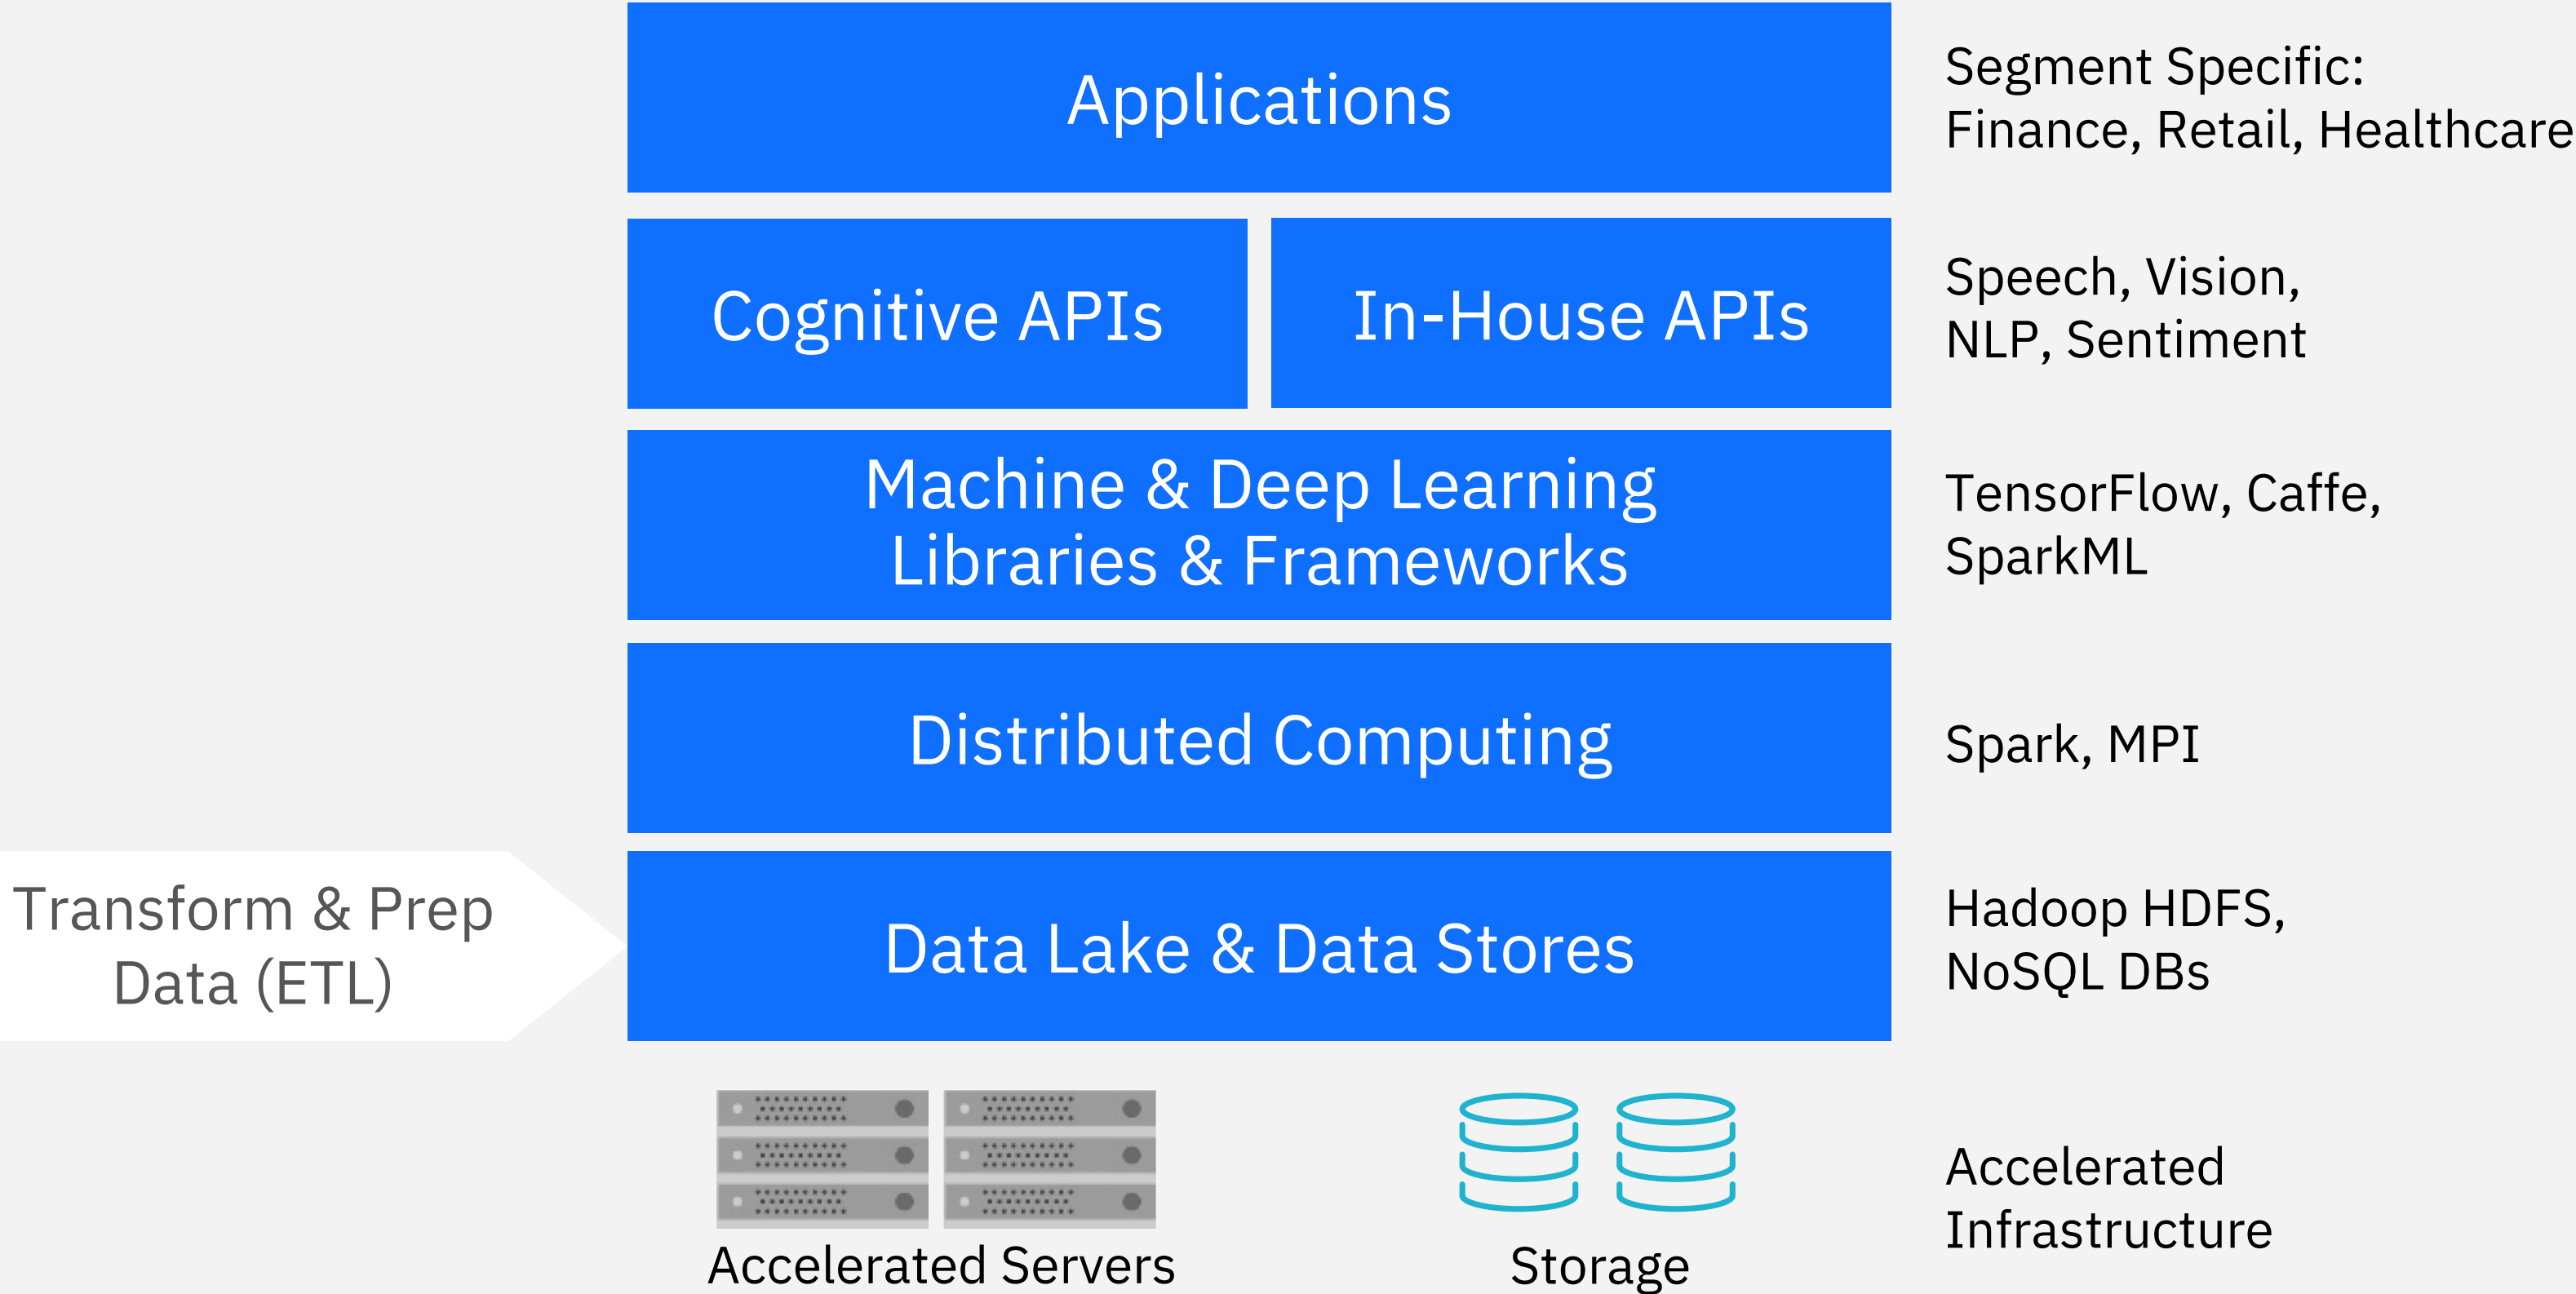
Conclude

Cluster Basics

It's a Process... and it all needs to work together in order to work correctly!

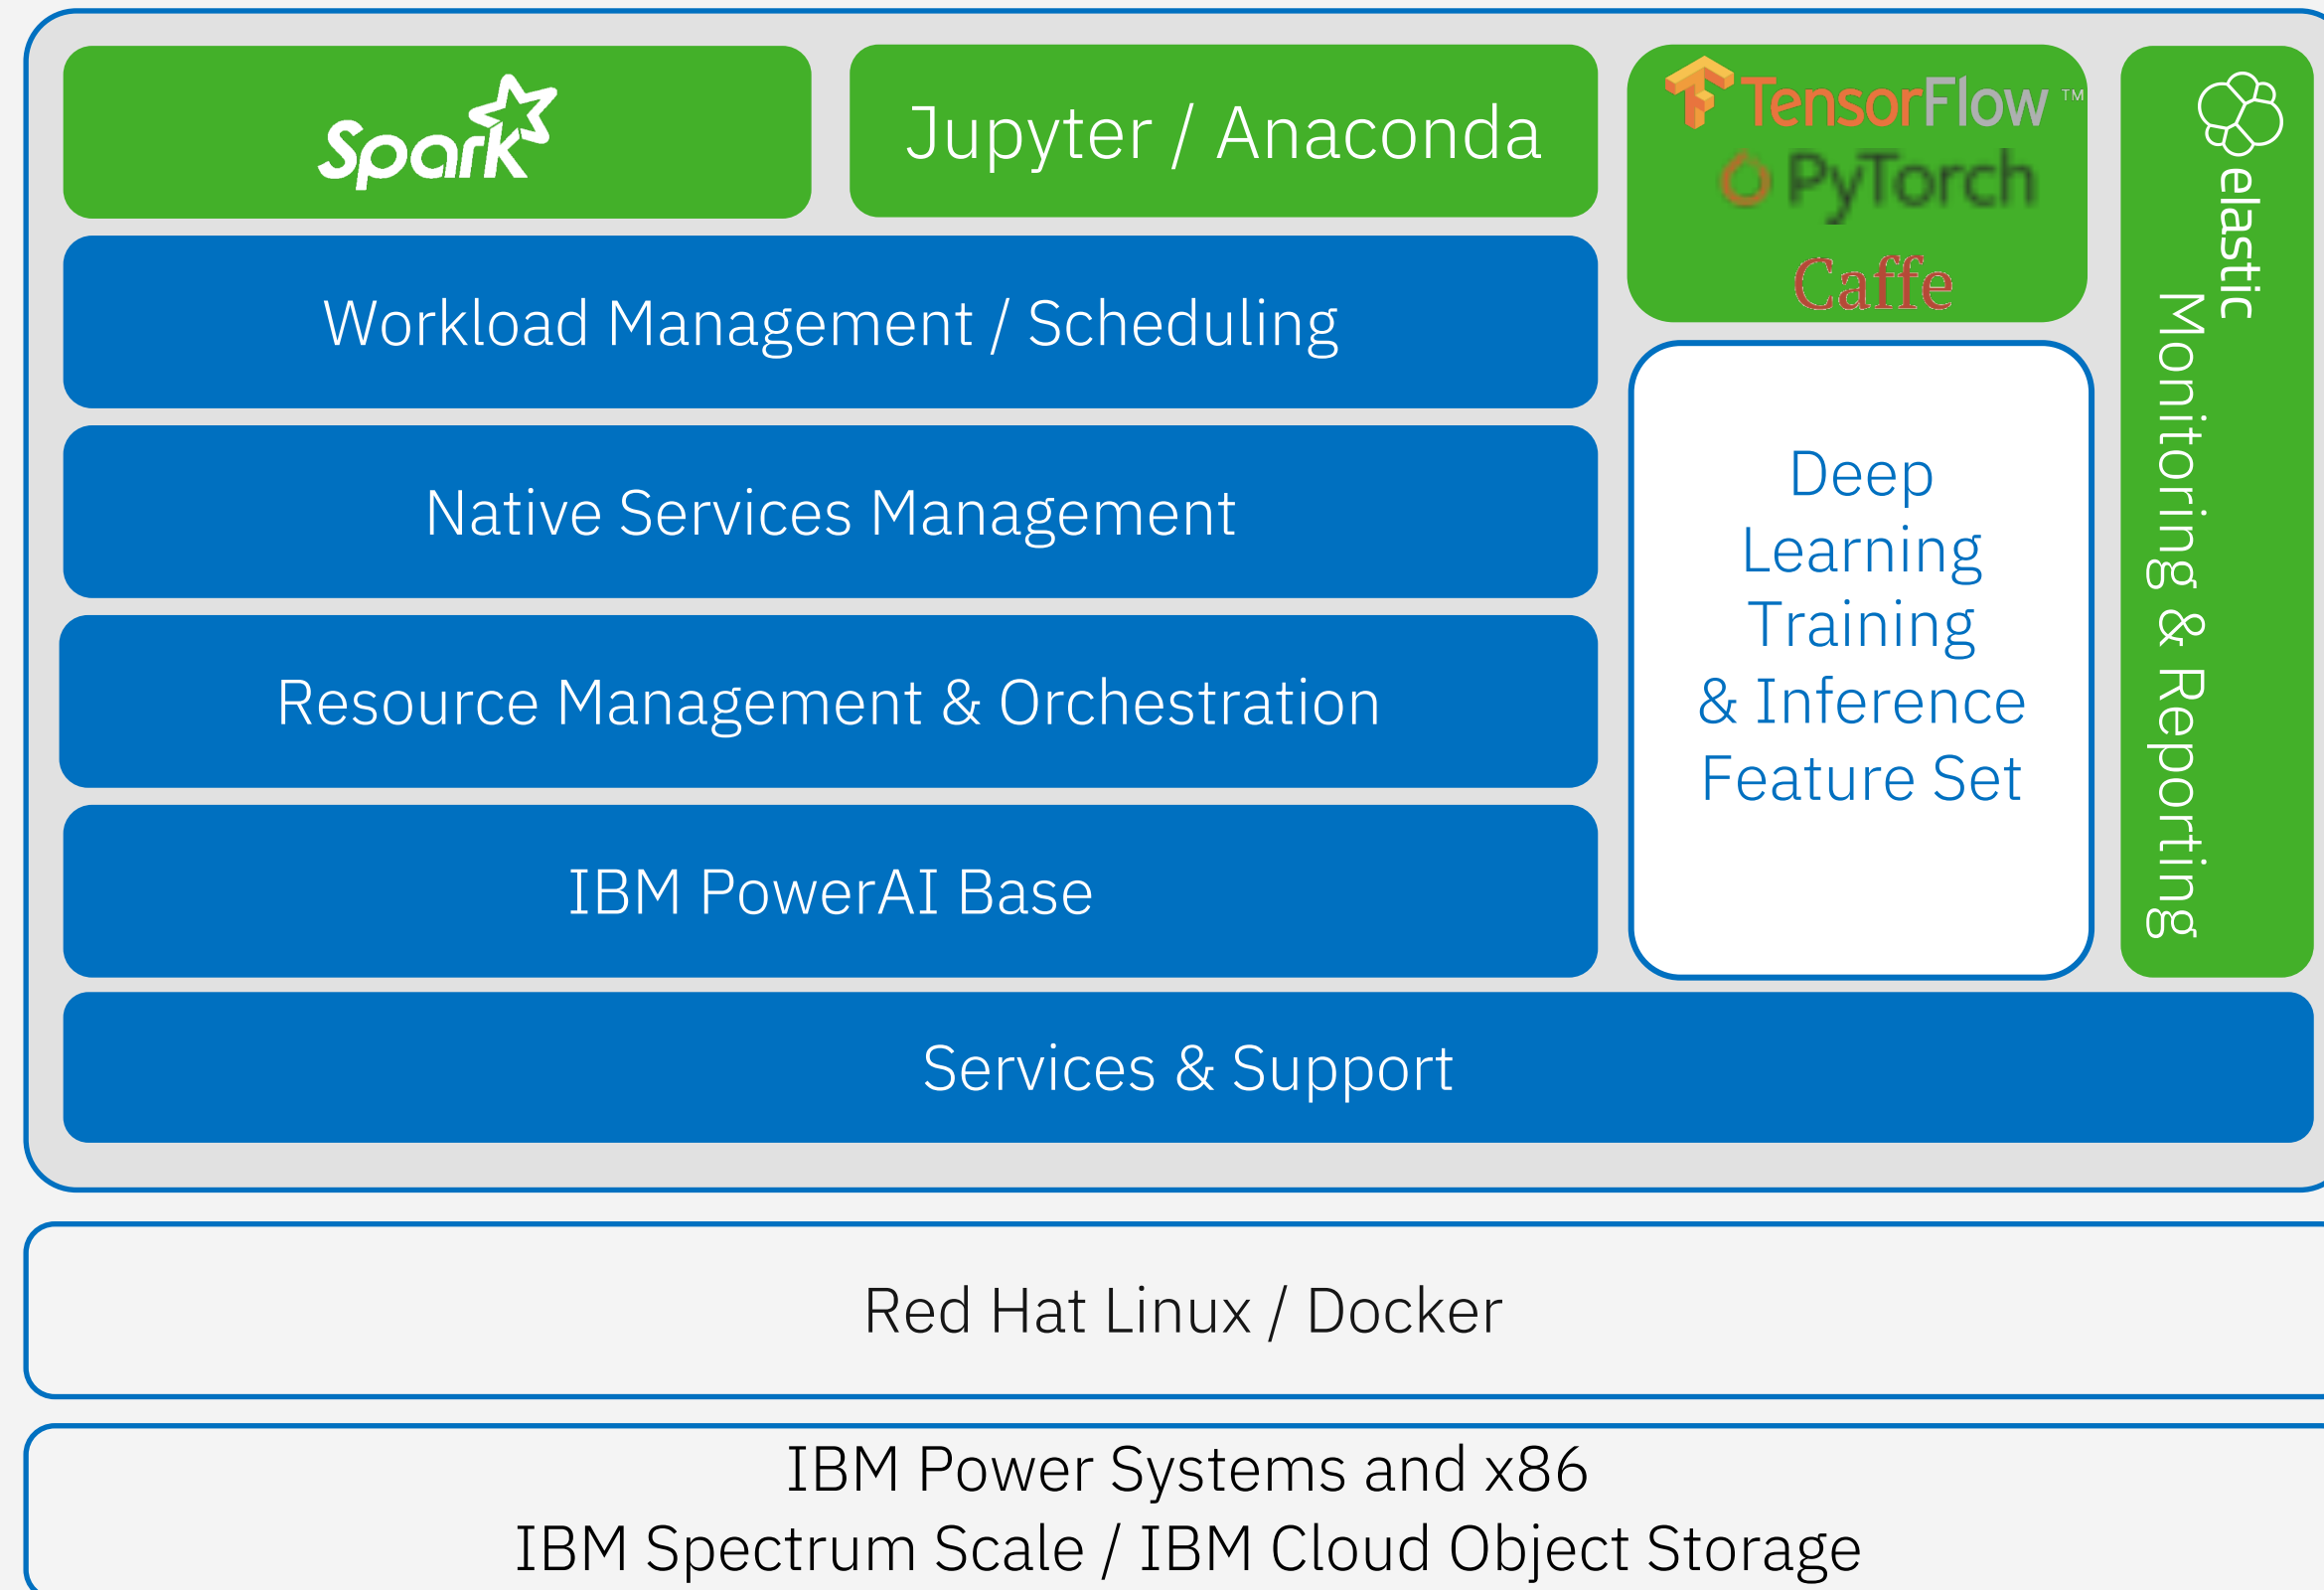


AI Infrastructure Stack



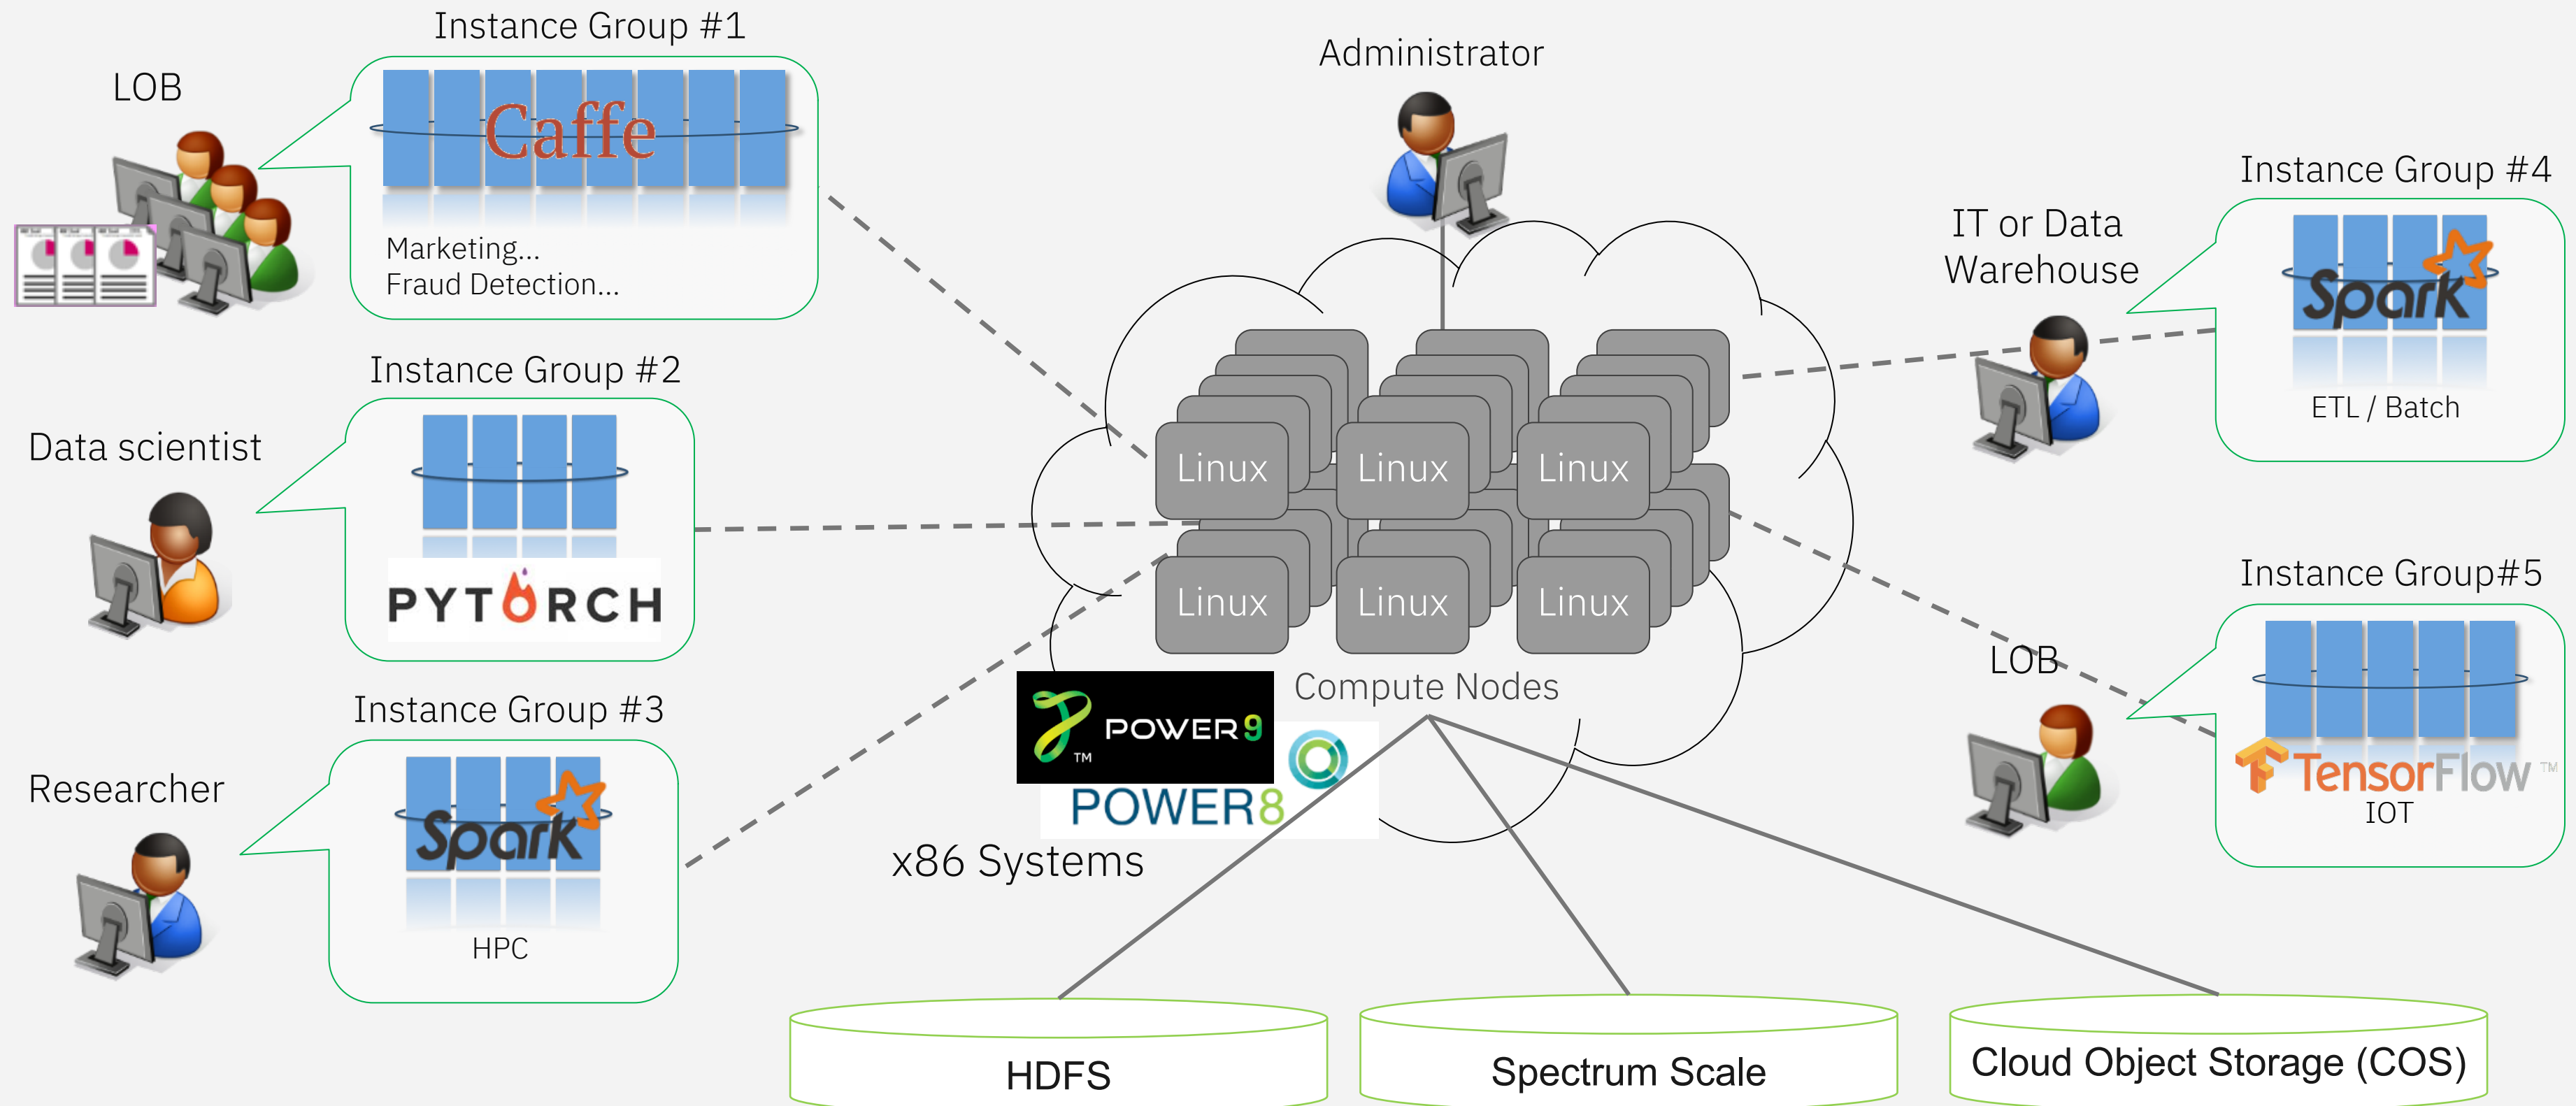
IBM Watson ML Accelerator:

Power AI Base + Spectrum Conductor + Deep Learning Impact



Multitenancy at Scale – Shared Services Architecture

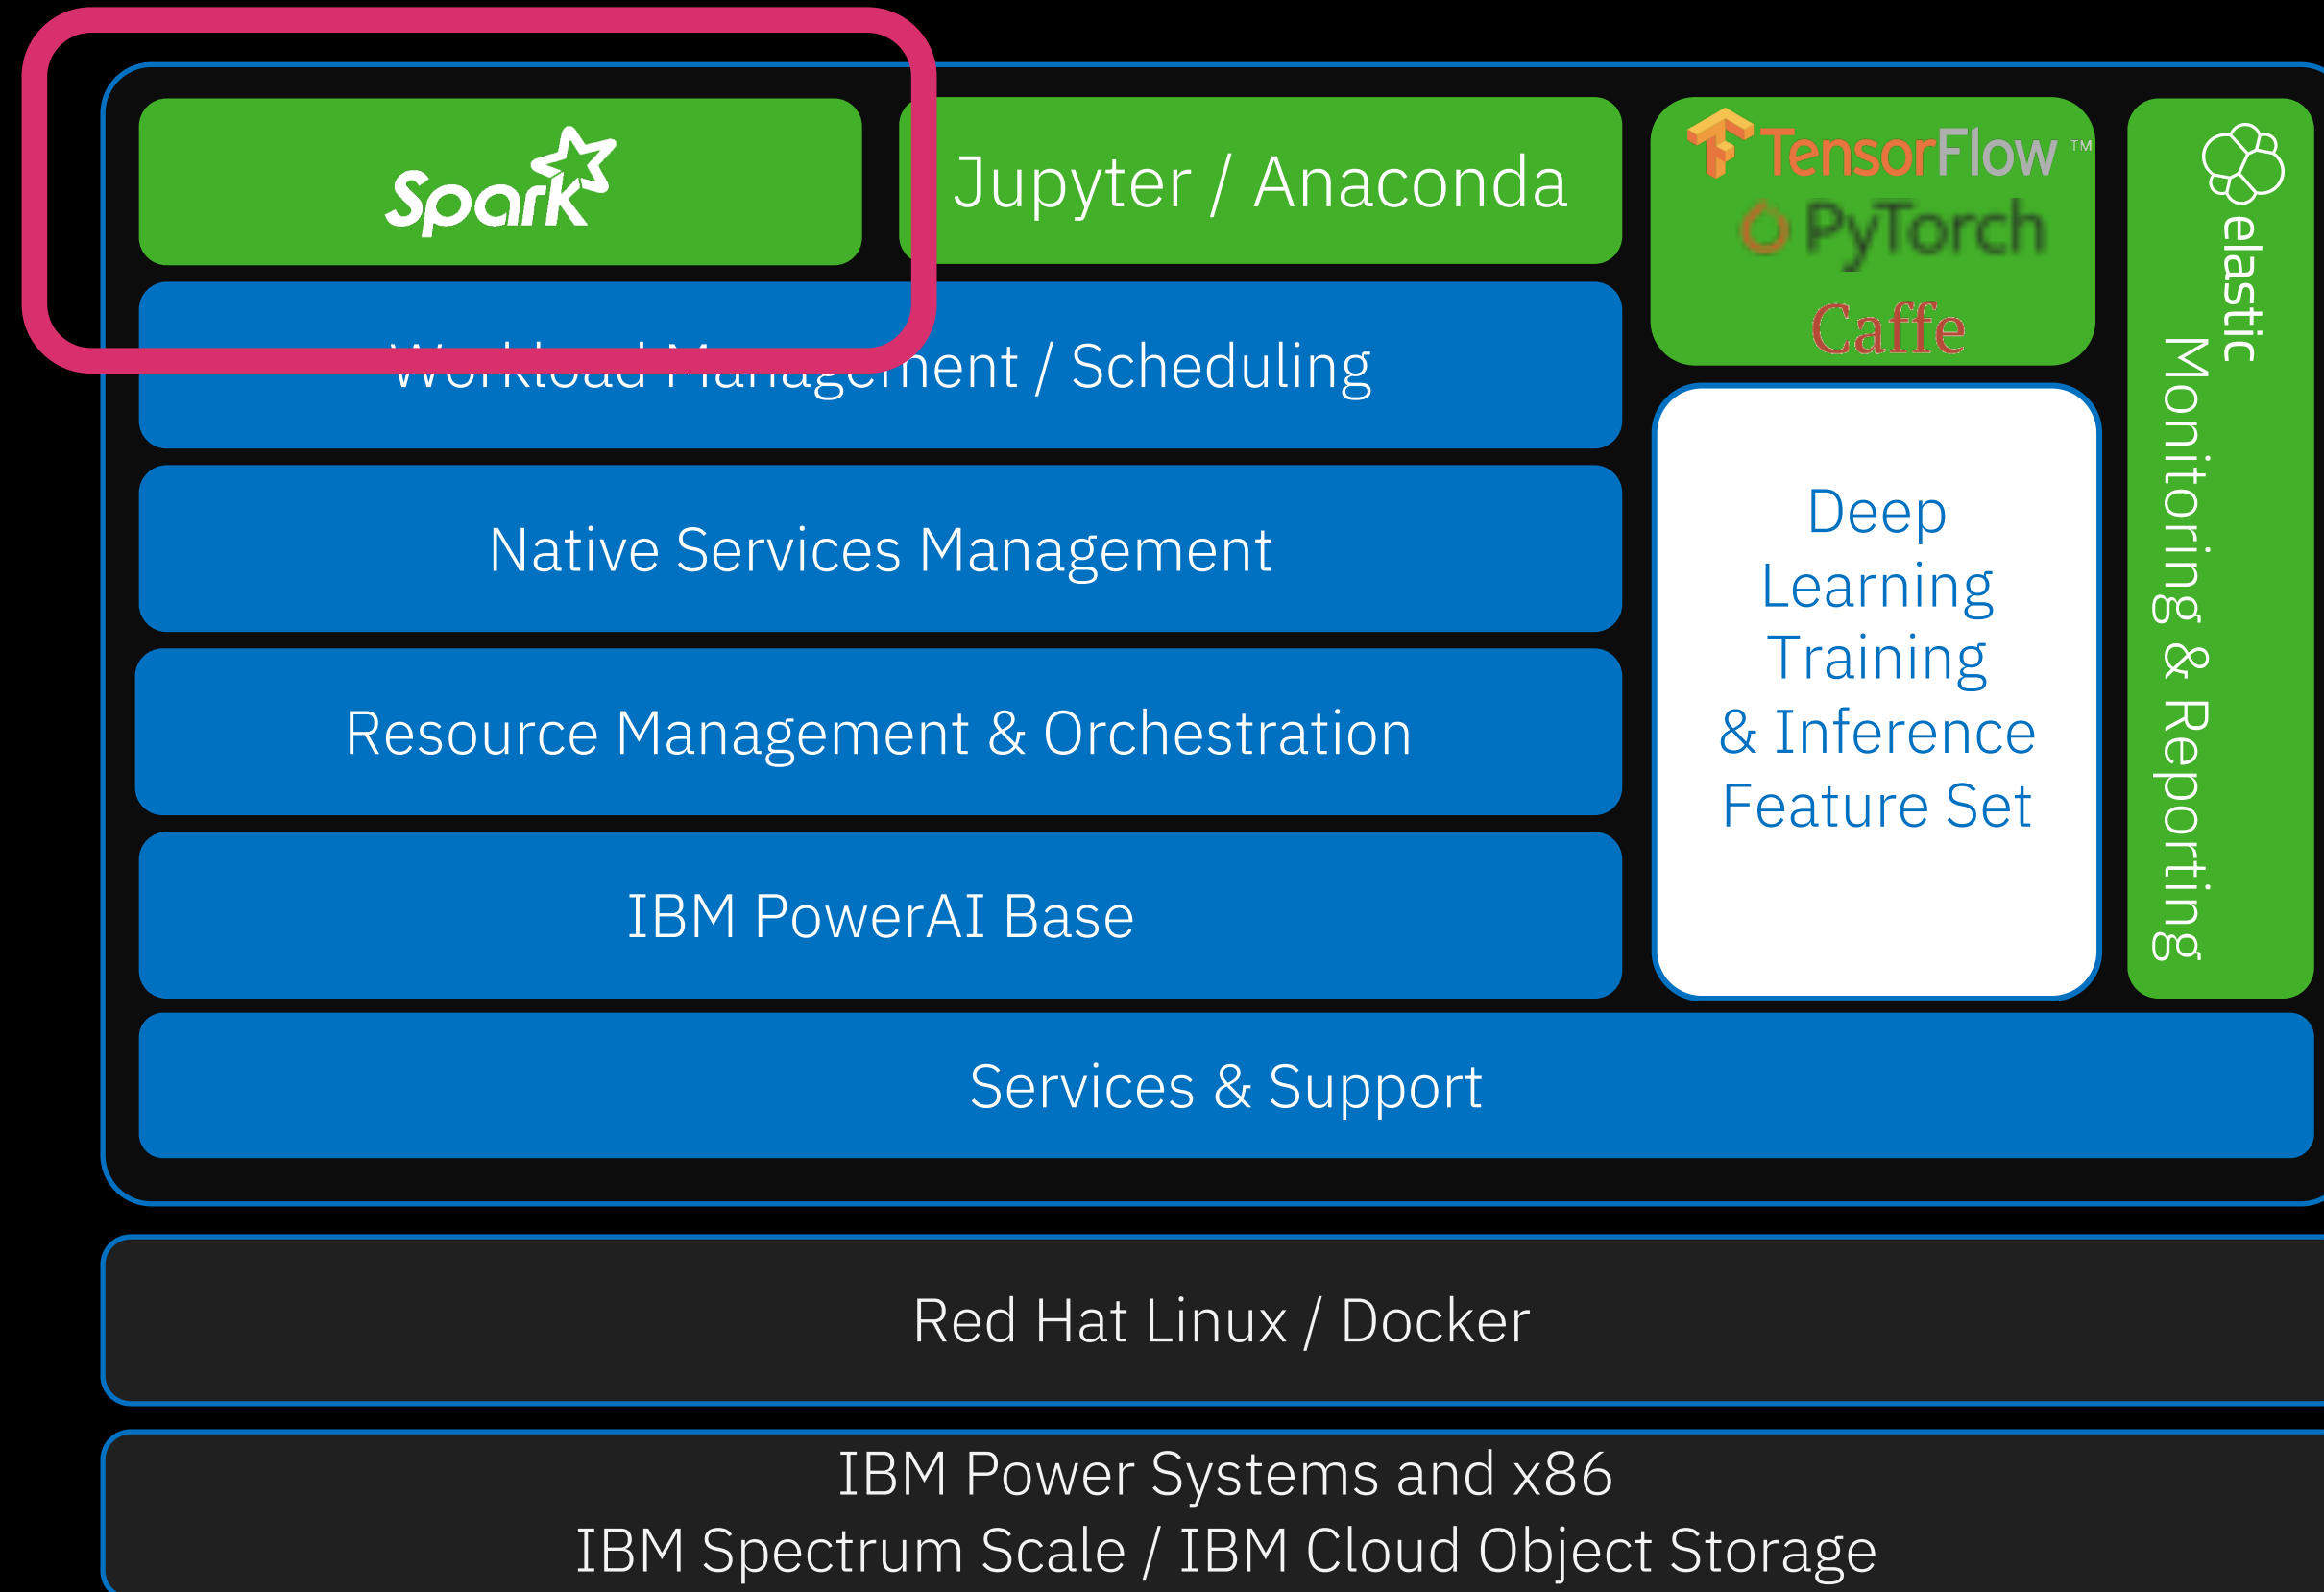
- **Physical view** Single cluster of resources managed by IBM **Watson Machine Learning Accelerator**
- **Logical view** Each group has **their own cluster resources** – Isolated & Secured with SLA Managed Resource Allocation



Cluster Access

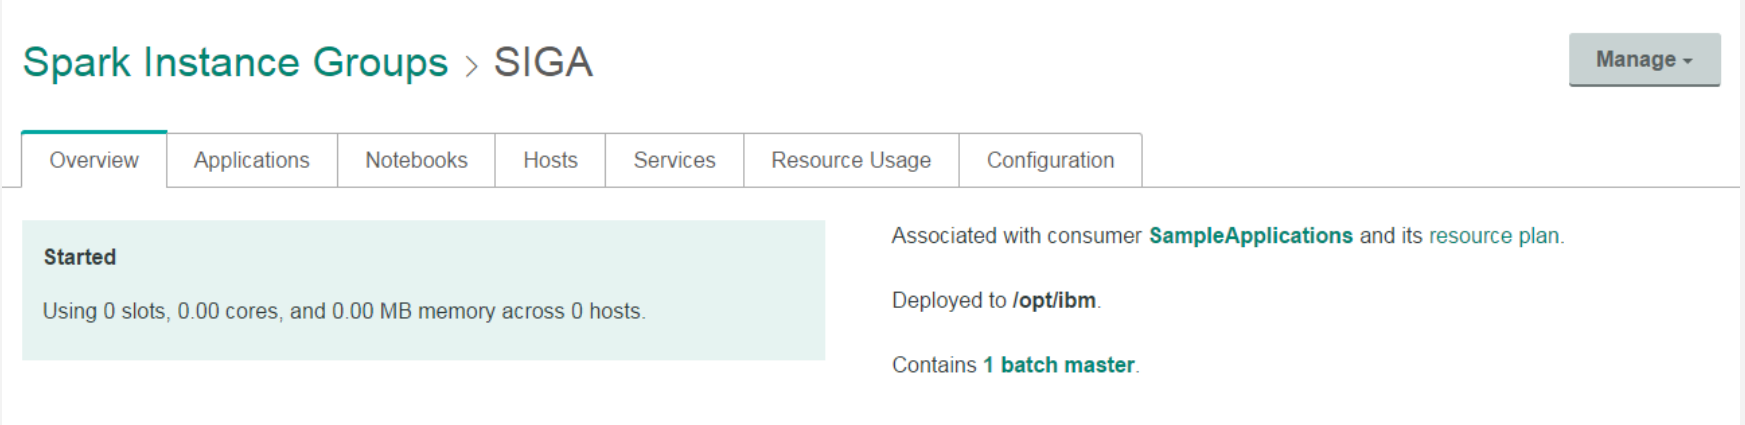
- URL <http://ibmminsky-head.cern.ch:8080/platform>
 - Use the fully qualified hostname (ibmminsky-head.cern.ch)
- User: [your cern.ch username](#)
 - Your username should be added to the appropriate cluster group by Admin team
- Password: [your cern.ch password](#)

Spark



Spark Instance Group Overview

- A [Spark instance group](#) is an installation of [Apache Spark](#) that can be scaled to run multiple masters, with history server, shuffle service, and notebooks configured.
- Each Spark instance group is a collection of services that is associated with a top-level consumer.
- You can manage and monitor a Spark instance group, and drill down to manage the related services and service instances.
- A Spark instance group can be [created to serve a line of business or a group member of a business organization](#).
- [Multiple versions](#) supported on the same cluster – (Spark 1.5.2, 1.6.0, 1.6.1, 2.0.1, 2.3.1, ...)



The screenshot shows the 'Spark Instance Groups' management interface for a group named 'SIGA'. At the top right is a 'Manage' button. Below the title is a navigation bar with tabs: Overview (selected), Applications, Notebooks, Hosts, Services, Resource Usage, and Configuration. The main content area displays the status 'Started' in a green box, indicating it is using 0 slots, 0.00 cores, and 0.00 MB memory across 0 hosts. To the right of this box, it states the group is associated with the 'SampleApplications' consumer and its resource plan, deployed to the '/opt/ibm' location, and contains 1 batch master.

Spark Instance Groups > SIGA Manage

Overview Applications Notebooks Hosts Services Resource Usage Configuration

Started
Using 0 slots, 0.00 cores, and 0.00 MB memory across 0 hosts.

Associated with consumer **SampleApplications** and its resource plan.
Deployed to **/opt/ibm**.
Contains **1 batch master**.

Spark Instance Group Overview

- **Runtime Isolation:**
 - Driver/executor process are owned by submitter
 - Within same cluster, users can deploy multiple Spark versions simultaneously on the same resources
- **Data at Rest Isolation:**
 - Data/log can only be accessed by owner
- **Management Isolation:**
 - RBAC authorization

The screenshot displays the 'Spark Instance Groups > SIGA' management page. At the top right is a 'Manage' button. Below the title is a navigation bar with tabs: Overview (selected), Applications, Notebooks, Hosts, Services, Resource Usage, and Configuration. The main content area shows the group's status as 'Started' in a teal box, with details: 'Using 0 slots, 0.00 cores, and 0.00 MB memory across 0 hosts.' To the right of this box, three informational lines are listed: 'Associated with consumer SampleApplications and its resource plan.', 'Deployed to /opt/ibm.', and 'Contains 1 batch master.'

Submitting Basic Spark Jobs

- Submitting workload via [Cluster management console](#)
 - Workload > Spark > Spark Instance Group > Submit Application
- Programmatically Using [REST API](#)

More Info:

Submitting batch job

https://www.ibm.com/support/knowledgecenter/en/SZU2E_2.3.0/tutorial/t_submit_app.html

Monitoring batch jobs

https://www.ibm.com/support/knowledgecenter/en/SZU2E_2.3.0/managing_applications/applications_monitoring_all.html

RESTful API

https://www.ibm.com/support/knowledgecenter/en/SZU2E_2.3.0/get_started/locating_rest_apis.html

YouTube Playlist

https://www.youtube.com/playlist?list=PLdgu63Noj6vLSTPPRdD7uuP9Zej_9stsO

Submitting Basic Spark Jobs

- Using spark-submit command from CLI

```
. /nfs/software/wmla/profile.platform
egosh user logon
cd spark-home-directory
bin/spark-submit --master spark://master-url:master-port --deploy-mode cluster --class
org.apache.spark.examples.SparkPi ./examples/jars/spark-examples_2.11-2.3.1.jar 100
```
- Find the master_url, port, and spark home directory in webui
(Workload > Spark > Spark Instance Group > Overview)

Spark... Even More

Shared RDDs / Dataframes

- Persistent Resilient Distributed Datasets (RDDs)
- Load once, read many, in native spark format

Adaptive Scheduling

- Run spark tasks on CPU and/or GPU based on scheduling logic

More Info:

Developing shared RDD applications

https://www.ibm.com/support/knowledgecenter/en/SSZU2E_2.3.0/managing_applications/applications_shared_develop.html

GPU RDDs

https://www.ibm.com/support/knowledgecenter/en/SSZU2E_2.3.0/developing_instances/instances_gpu_run_app_gpurd.html

GPU RDD example

https://www.ibm.com/support/knowledgecenter/en/SSZU2E_2.3.0/reference_s/gpu_rdd_samples.html

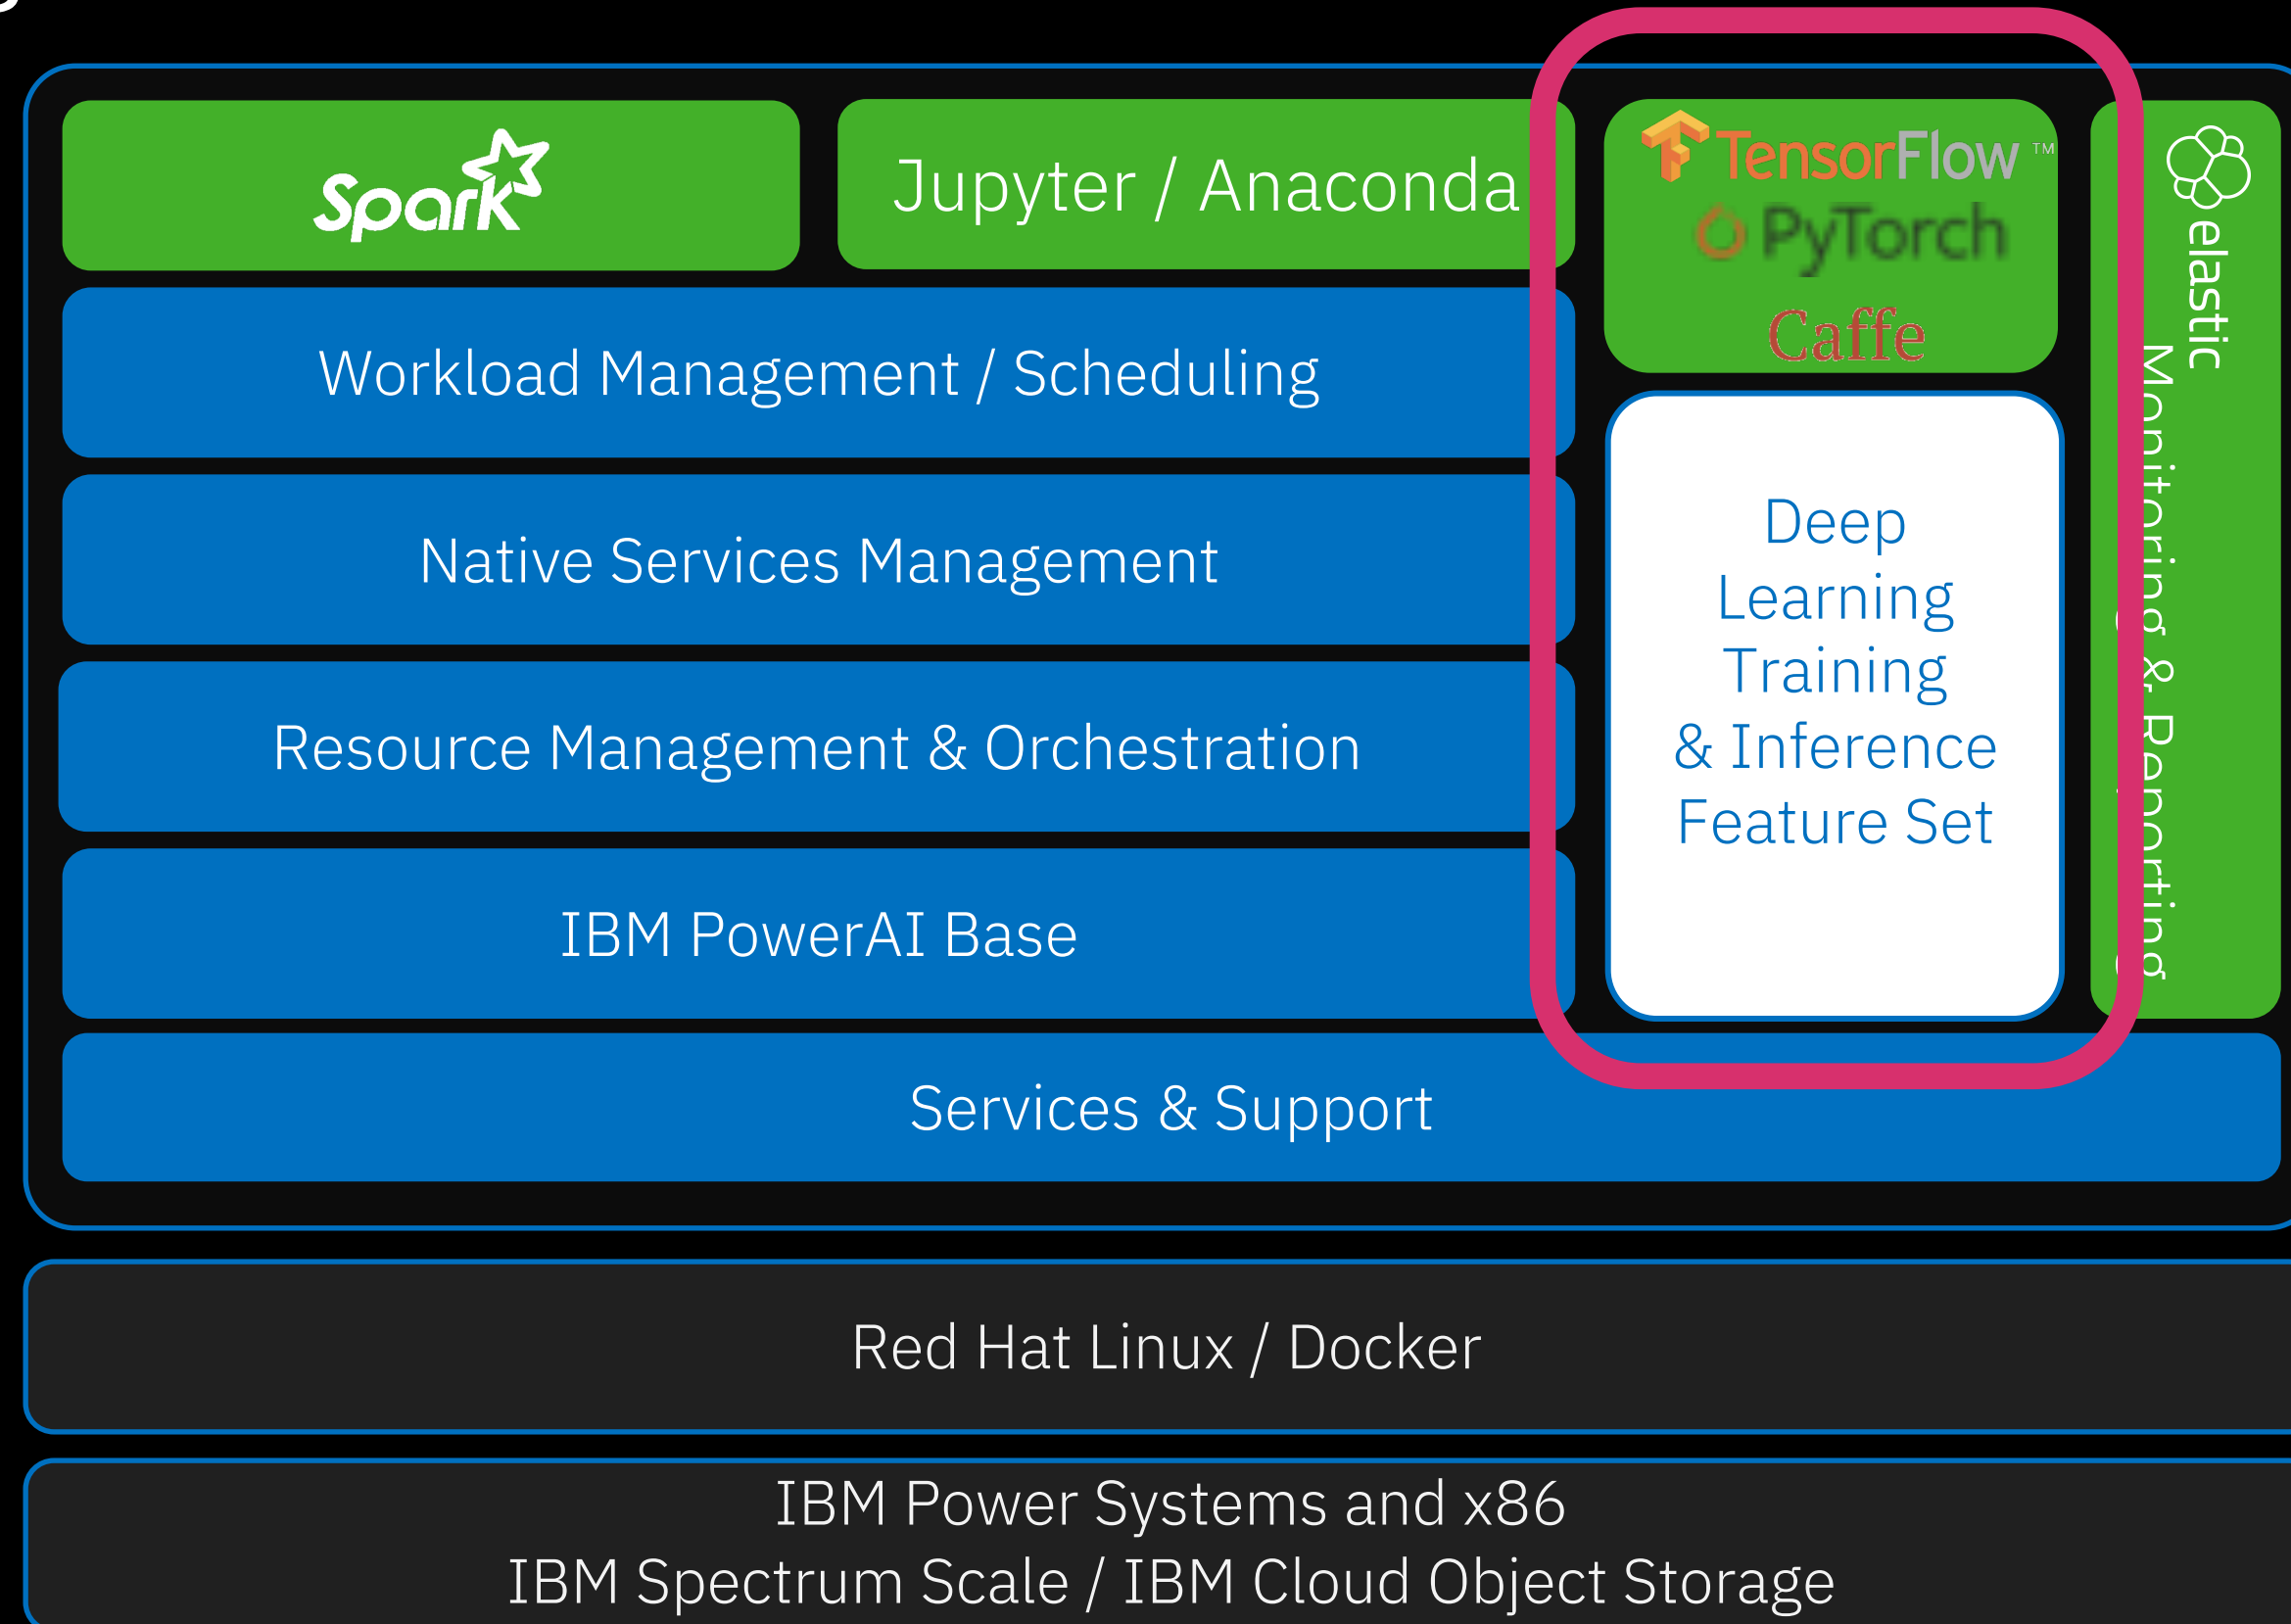
Guidelines for application code

https://www.ibm.com/support/knowledgecenter/en/SSZU2E_2.3.0/performance_tuning/application_code_considerations.html

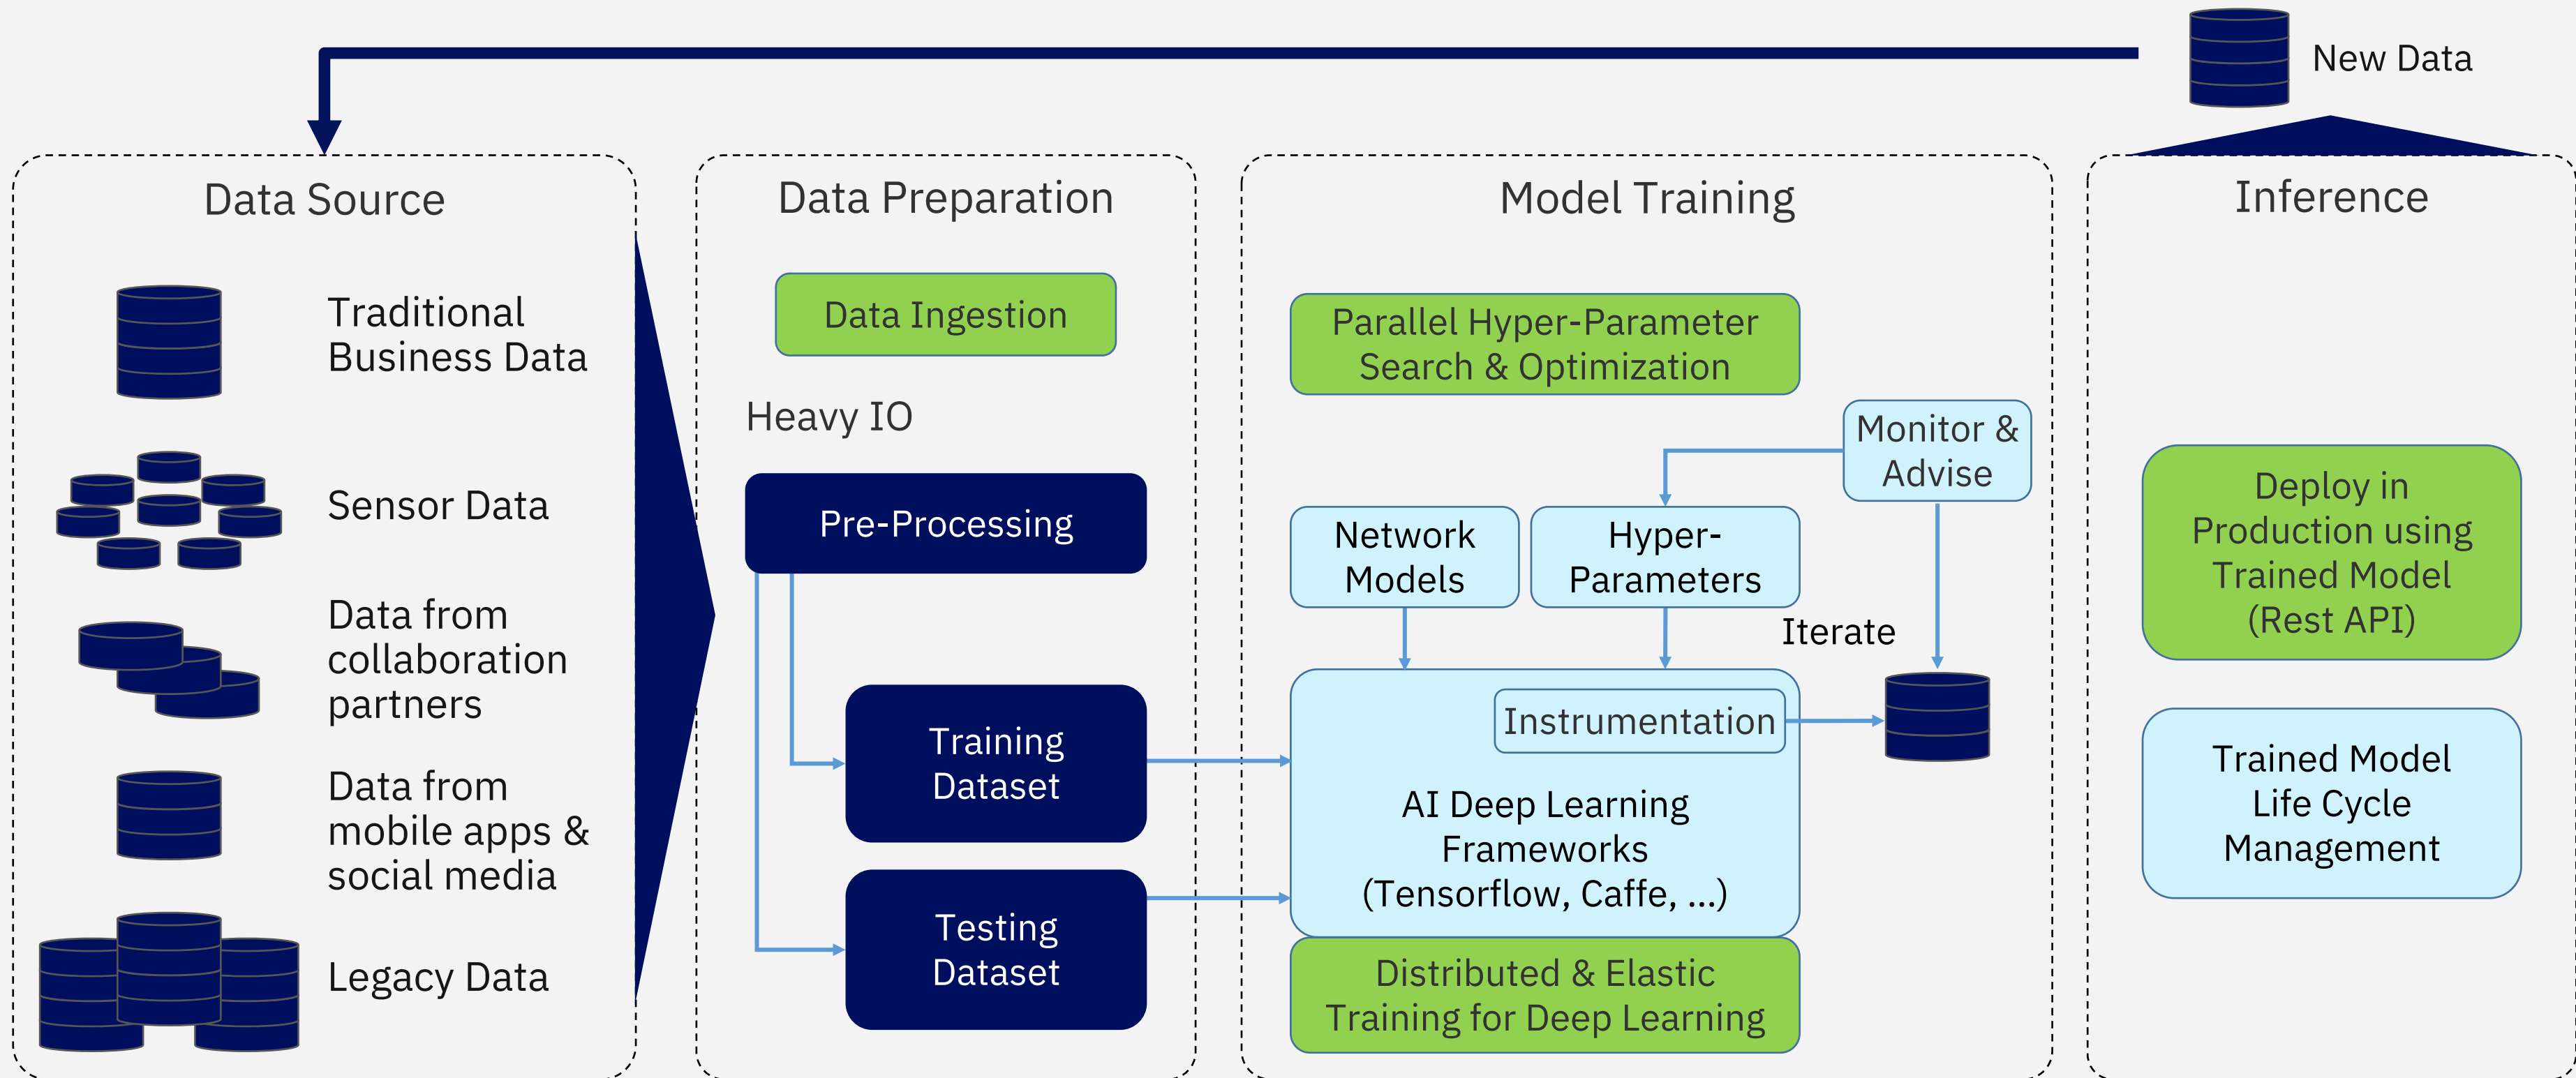
Adaptive Scheduling

https://www.ibm.com/support/knowledgecenter/en/SSZU2E_2.3.0/developing_instances/instances_gpu_adaptive_scheduling.html

Deep Learning



Basic Workflow for ML/DL Training & Execution



Importing Datasets

- Create different types of datasets
 - LMDB
 - TFRecords
 - Images for object classification
 - Images for object detection
 - Images for vector output
 - CSV files
 - Others (COPY, NLP (NER, POS, Segmentation), Text classification)

More Info:

https://www.ibm.com/support/knowledgecenter/SFHA8_1.2.0/us/deep-learning-create-dataset.html

Running Basic Training Jobs

- Importing a model from model zoo
- Creating a new model from template
- Launch a training job
- Launch hyperparameter tuning

More Info:

Create a training model

https://www.ibm.com/support/knowledgecenter/SSFHA8_1.2.0/us/deep-learning-create-training-model.html

Edit models for training:

https://www.ibm.com/support/knowledgecenter/SSFHA8_1.2.0/us/configure-caffe-model.html

https://www.ibm.com/support/knowledgecenter/SSFHA8_1.2.0/us/configure-tensorflow-model.html

https://www.ibm.com/support/knowledgecenter/SSFHA8_1.2.0/us/configure-pytorch-model.html

Model zoo:

<https://us-south.git.cloud.ibm.com/ibmconductor-deep-learning-impact>

End to end workflow:

<https://developer.ibm.com/tutorials/use-computer-vision-with-dli-watson-machine-learning-accelerator/>

Hyperparameter Tuning

- Find optimal hyperparameters faster
 - Search algorithms (random, TPE, Bayesian, Hyperband)
 - Parallel search maximizing resource use
 - Option for custom parameters

More Info:

Edit model for hyperparameter tuning

https://www.ibm.com/support/knowledgecenter/SFHA8_1.2.0/us/configure-tensorflow-model-hyperparameter-tuning.html

https://www.ibm.com/support/knowledgecenter/SFHA8_1.2.0/us/configure-pytorch-model-hyperparameter-tuning.html

Elastic Distributed Training Example

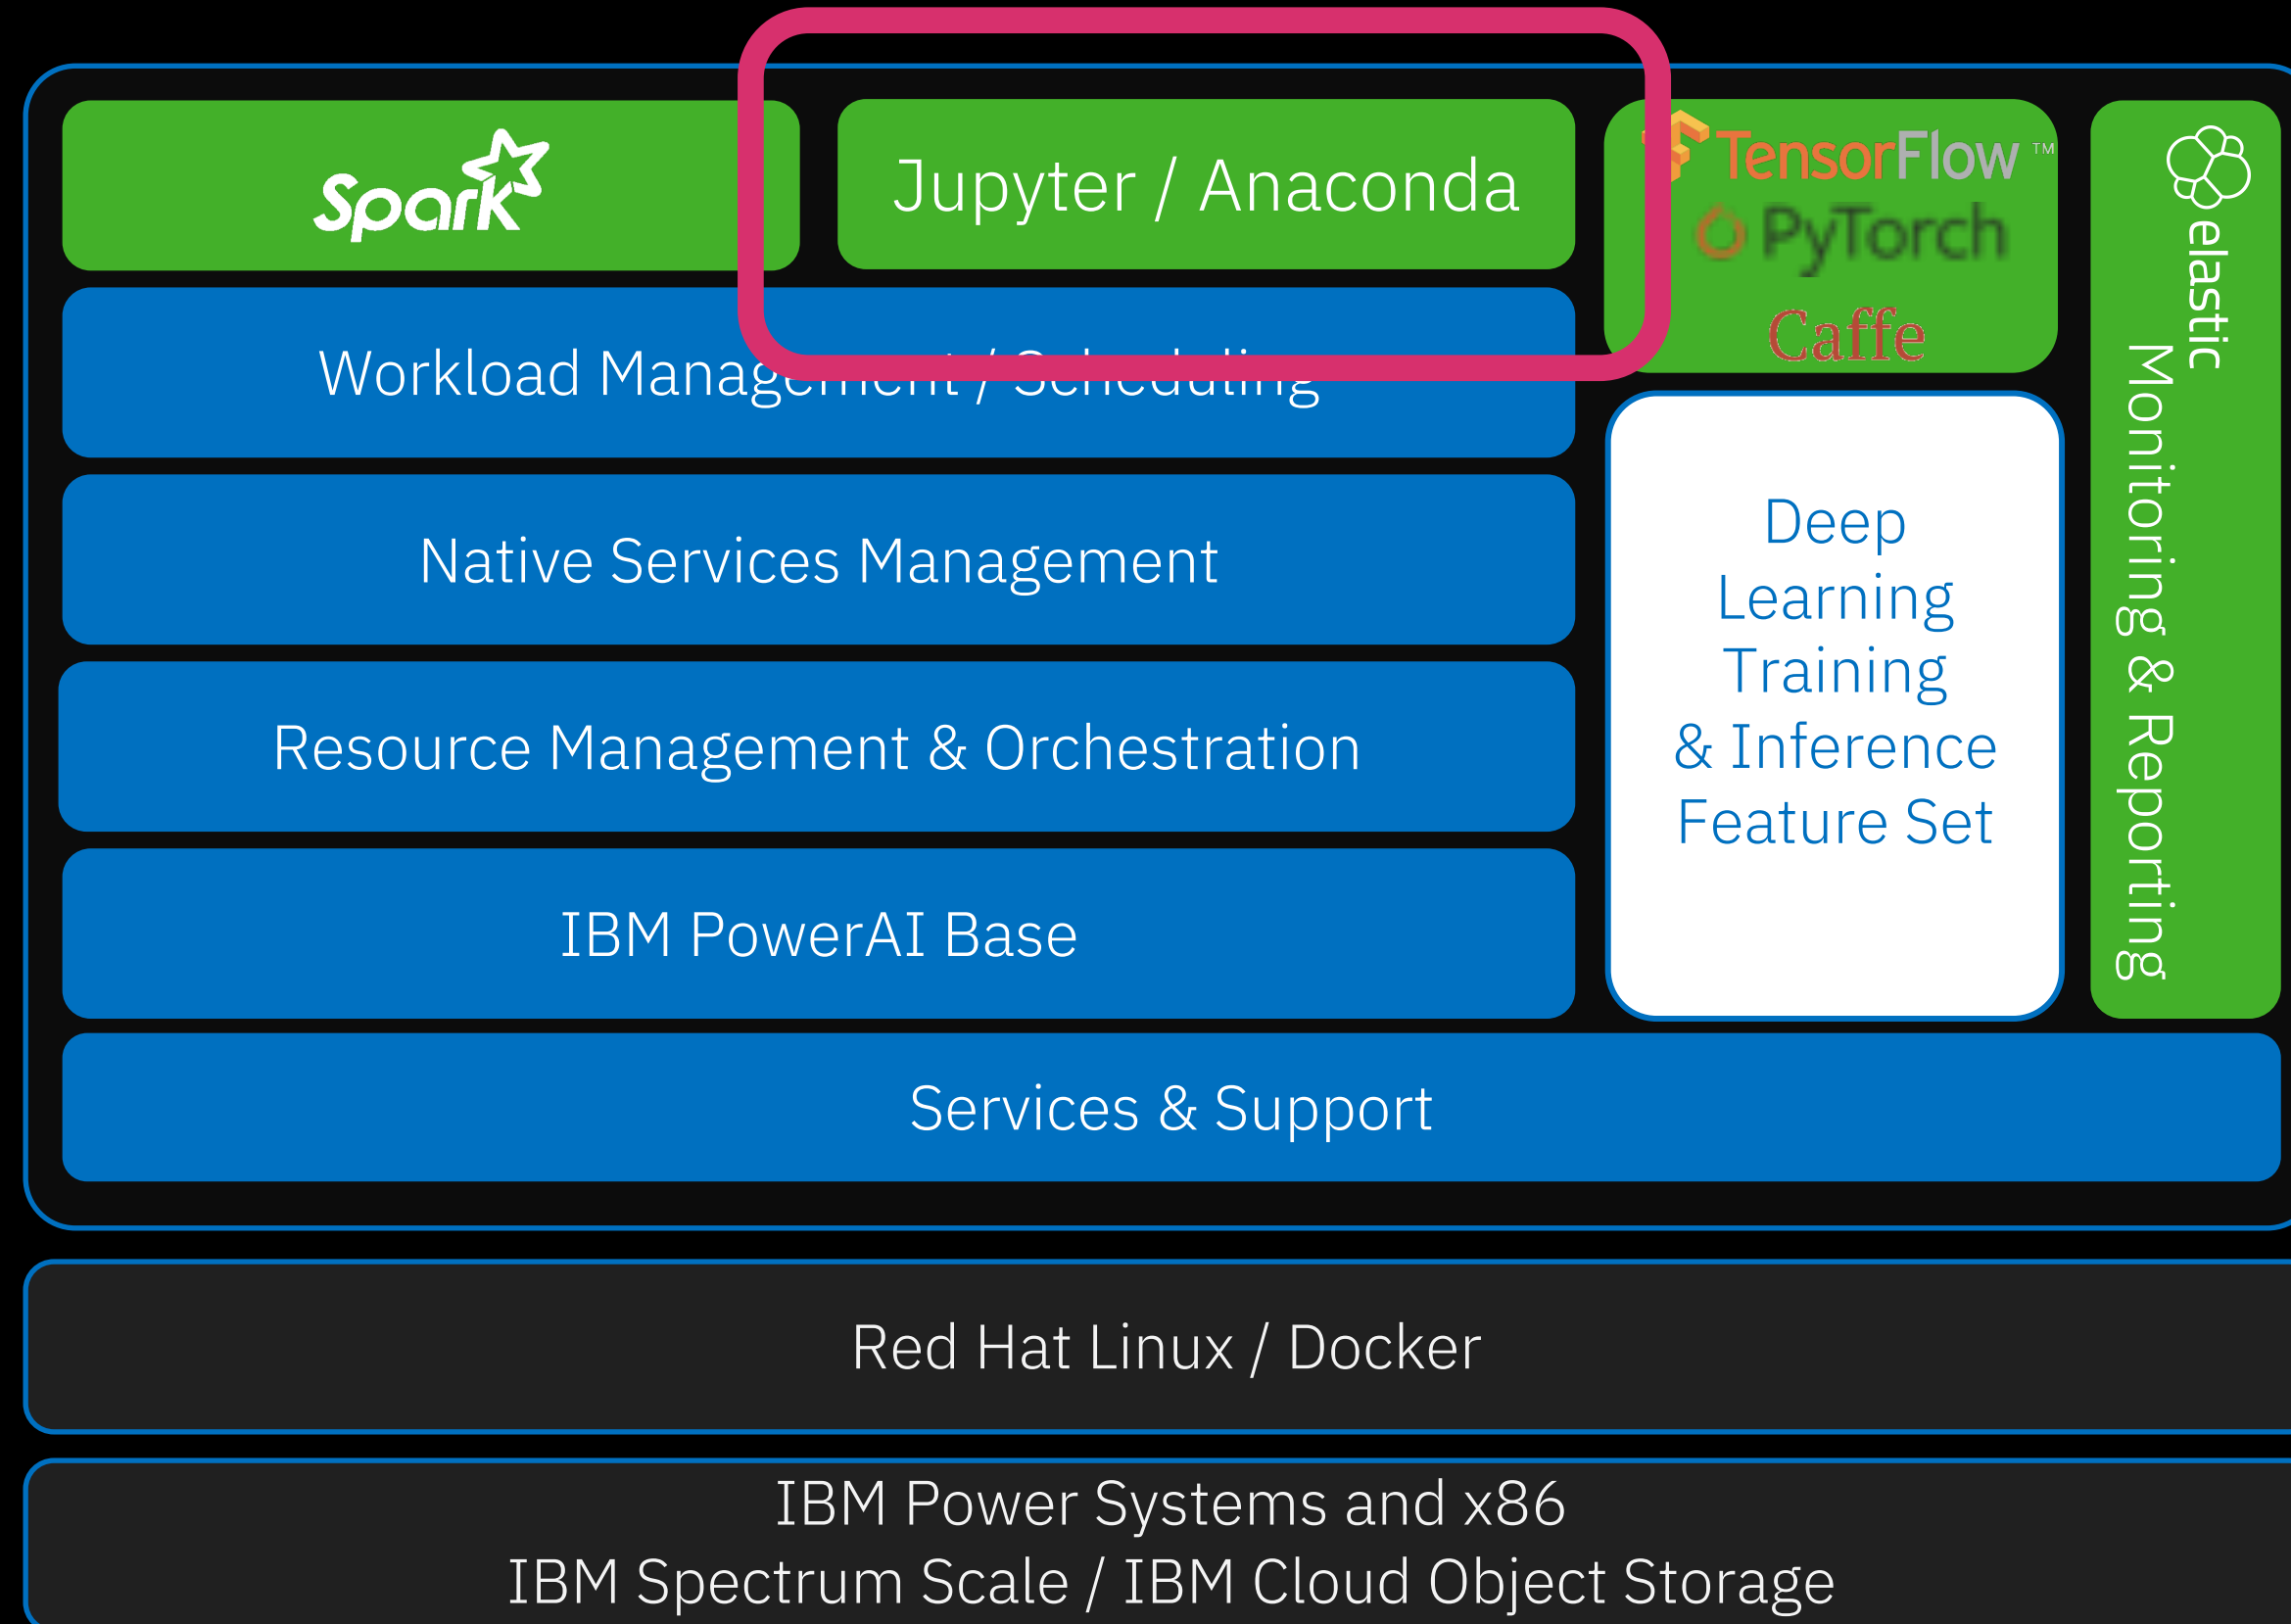
- Single-node, multi-GPU training support without parameter server .
- [Multi-node](#), [multi-GPU](#) training support with parameter server.
- Both [synchronous](#) and [asynchronous](#) distributed training algorithms support for multi-node, including synchronous gradient data control algorithm, and asynchronous gradient data control algorithm.
- NVIDIA Collective Communications Library (NCCL) support for broadcasting, reducing gradient and weight data across multiple GPUs.
- Review modifications for model for EDT

More Info:

Edit model for elastic distributed training

https://www.ibm.com/support/knowledgecenter/SFHA8_1.2.0/us/configure-pytorch-model-elastic-distributed-training.html

Jupyter / Anaconda



Jupyter / Anaconda Overview

- Distributed anaconda distribution, conda environment, and conda package management
- Option to link to Jupyter notebook
 - Jupyter notebook linked to underlying conda environment
 - Conda environment managed per group and ability to create multiple
 - Collaboration option
 - Easy interface into cluster (development, data management / cleansing, etc.)

More Info:

Notebooks (snapml, xgboost, keras, and more)

<https://github.com/IBM/wmla-assets>

Tutorials

<https://developer.ibm.com/tutorials/category/watson-ml-accelerator/>

Open a notebook and create a note

https://www.ibm.com/support/knowledgecenter/en/SSZU2E_2.3.0/tutorial/t_notebook.html

Work with anaconda through notebook

- Workload > Spark > My Application and Notebooks > Open Notebook
- If you don't see this, maybe your users needs a notebook assigned to it or collaboration enabled

More Info:

Notebooks (snapml, xgboost, keras, and more)

<https://github.com/IBM/wmla-assets>

Tutorials

<https://developer.ibm.com/tutorials/category/watson-ml-accelerator/>

Open a notebook and create a note

https://www.ibm.com/support/knowledgecenter/en/SSZU2E_2.3.0/tutorial/t_notebook.html

Work with anaconda through CLI

- Activate the Anaconda Distribution instance

```
source  
/nfs/public/wmla/anaconda/[group]/201903/anaco  
nda/etc/profile.d/conda.sh
```

- Activate the conda environment

```
conda activate dlipy3
```

More Info:

Notebooks (snapml, xgboost, keras, and more)

<https://github.com/IBM/wmla-assets>

Tutorials

<https://developer.ibm.com/tutorials/category/watson-ml-accelerator/>

Open a notebook and create a note

https://www.ibm.com/support/knowledgecenter/en/SSZU2E_2.3.0/tutorial/t_notebook.html

Open Discussion / Questions

Thank you!!

Eric Fiala
Cognitive Solution Architect

—

Eric.J.Fiala@ibm.com
ibm.com

Continue the discussion on Slack
[#cern-openlab](https://ibm-systems-cern.slack.com)

